IPA AI Usage Guidelines

Introduction

Artificial intelligence tools offer tremendous potential to advance IPA's mission of reducing poverty through evidence-based research. AI can help us write more effectively, analyze data more efficiently, conduct research more thoroughly, and collaborate more productively. As a data collection and research organization committed to both innovation and responsible stewardship, we want to enable staff to use AI tools effectively while protecting the sensitive information entrusted to us by research participants, donors, and partners.

These guidelines are designed to help you:

- Understand when you can use AI tools without seeking approval
- Determine what types of information are safe to share with AI tools
- Know when and how to request approval for new use cases
- Use AI responsibly in your research and operational work

Our approach balances innovation with protection. Most of your AI use won't require special approval - the key is understanding how to classify your data and choosing appropriate tools.

AI Policymaking at IPA AI technology is rapidly evolving, so this document will be updated periodically. It is maintained by IPA's Management and Information Systems Technology (MIST) team under the supervision of the Senior Management Team, in close collaboration with experts from the Research, Operations, and Legal departments and with input from users throughout IPA. MIST serves as the organizing body that will receive all inquiries and route them to Research, MIST, Legal or other teams as necessary for advice and approvals.

START HERE: Classify Your Data

Before using any AI tool, you must first classify the data you plan to work with.

IPA uses a four-tier data classification system based on the impact that unauthorized disclosure would have on IPA, research participants, donors, or partners.

Data Classification Table

Classification	Definition	Examples	Can I Use AI?
PUBLIC	Information produced for publication with no impact from disclosure	Published research papers, marketing materials, public blog posts, IPA website content, published anonymized datasets	✅ Yes - any AI tool
INTERNAL	Non-sensitive organizational data; disclosure would cause minor, recoverable inconvenience	Project documentation, non-public training materials, internal meeting notes (non-confidential), draft documents without sensitive content	✅ Yes - with IPA-approved tools only
CONFIDENTIAL	Sensitive data; disclosure could cause significant harm to individuals or IPA	Datasets with PII (names, emails, locations), PII of any third party or staff, unpublished research data, embargoed research findings, grant proposals in development, salary information, donor correspondence, IRB protocols, information covered by NDAs or confidentiality agreements/terms in awards, proprietary information from donors or partners, information restricted by data protection laws	⚠️ Limited - approved tools only OR requires approval
HIGHLY CONFIDENTIAL	Most sensitive data; disclosure could cause severe harm, legal liability, or reputational damage	Datasets with sensitive personal information (health data, biometric data, sexual orientation, political/religious views,children's data) or anyone's sensitive data (including staff or third parties), highly confidential business information subject to legal restrictions, information that could endanger research participants or violate consent, merger/acquisition details, material covered by strict contractual confidentiality obligations	🛑 Contact MIST first - requires strict controls

Decision Tree

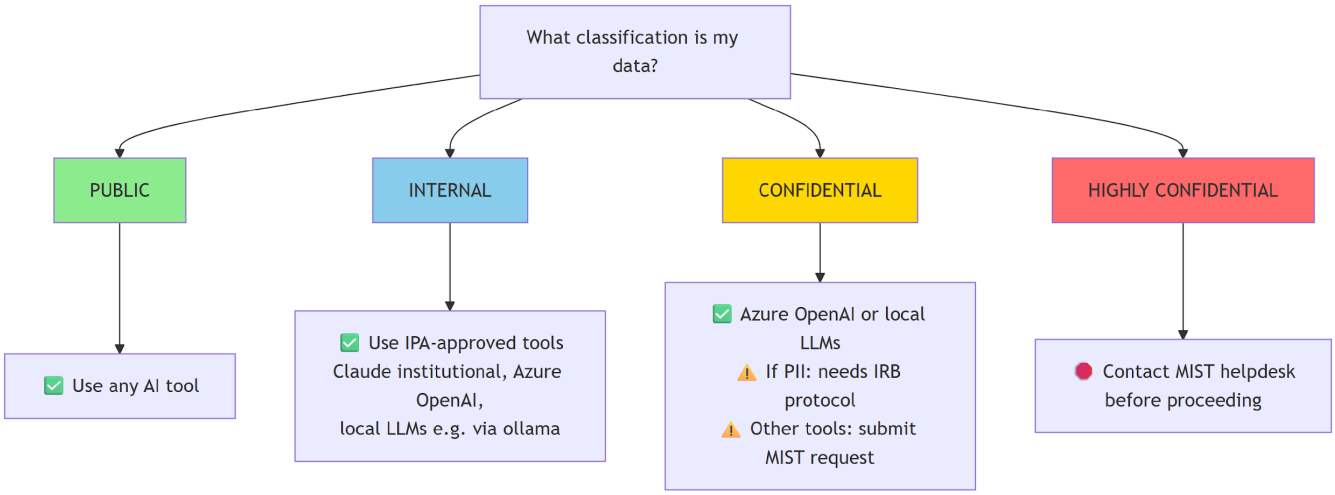

When in doubt about classification, choose the more restrictive category and/or contact MIST.

Understanding Data Classification: Key Distinctions

PII vs. Sensitive Personal Information (SPI)

The distinction between CONFIDENTIAL and HIGHLY CONFIDENTIAL data often hinges on the type of personal information and the level of sensitivity of documents (e.g., grant agreements with strategic information may be HIGHLY CONFIDENTIAL even without PII or licensed information that IPA is not authorized to share).:

Regular PII (CONFIDENTIAL):

- Contact information: names, email addresses, phone numbers, mailing addresses
- Standard identifiers: participant IDs that could be linked back to individuals, employee numbers
- Demographic information: age, gender, occupation, education level
- Location data: specific addresses, GPS coordinates
- Financial information: income levels, transaction data

Sensitive Personal Information - SPI (HIGHLY CONFIDENTIAL):

- Health and medical data (diagnoses, treatments, HIV status, disabilities)
- Biometric data (fingerprints, facial recognition, DNA, iris scans)
- Sexual orientation or gender identity

- Political opinions, religious or philosophical beliefs
- Trade union membership
- Criminal history, arrests, or allegations
- Data about children under 18 (especially under 13)
- Information that could endanger individuals if disclosed

Critical rule: If your dataset contains ANY SPI, the entire dataset is classified as HIGHLY CONFIDENTIAL, even if most fields are regular PII.

De-identification Standards

What counts as "properly de-identified"?

De-identified data has had all direct and indirect identifiers removed or modified so that individuals cannot be reasonably identified. However, de-identification exists on a spectrum:

Minimally acceptable de-identification (still CONFIDENTIAL):

- Direct identifiers removed (names, contact info, ID numbers)
- BUT: Combination of remaining variables could still identify individuals
- Example: Age (45) + Occupation (village chief) + District (specific) + Gender = potentially identifiable in small communities

Strong de-identification (may qualify as INTERNAL if no other restrictions apply):

- All direct identifiers removed
- Indirect identifiers generalized or removed
    - Ages grouped into ranges (25-34, 35-44)
    - Geographic data limited to region, not specific village
    - Rare occupations generalized or suppressed
    - Dates rounded to month or quarter
- K-anonymity achieved: each individual is indistinguishable from at least k-1 others
- No outliers or unique combinations that could identify individuals

Important: Even strongly de-identified data may still be CONFIDENTIAL if:

- It's unpublished research data subject to embargo
- Covered by donor agreements restricting data sharing
- Governed by IRB protocols that prohibit third-party sharing
- Contains information that could harm IPA if disclosed (e.g., negative findings about a partner)

Indirect Identifiers and Re-identification Risk

What are indirect identifiers?

Variables that alone don't identify someone, but in combination can:

Common indirect identifiers in IPA research:

- Geographic information: village, district, neighborhood, school name
- Demographic combinations: age + gender + occupation + household size
- Rare characteristics: very high/low income, unusual occupation, unique circumstances
- Temporal information: exact dates of participation, birth dates
- Participation in specific programs: if program served few people

Examples of re-identification risk:

> ⚠ Example: Small Village Survey
>
> Dataset contains: Age (52), Gender (Female), Occupation (Teacher), Village (Kakuma)
>
> Even without names, if Kakuma village has only one 52-year-old female teacher, this person is identifiable. This data requires CONFIDENTIAL classification at minimum.

> ⚠ Example: Health Intervention Study
>
> Dataset contains: Age range (45-54), District (North Region), Received HIV testing (Yes), Has diabetes (Yes)
>
> Even with age ranges, the combination of health conditions might identify individuals in a small sample. This is HIGHLY CONFIDENTIAL due to SPI (health data).

How to assess re-identification risk:

1. **Population size:** Smaller populations = higher risk
   - Survey of 5 village chiefs in a district → HIGH RISK even if anonymized
   - National representative survey of 10,000 households → LOWER RISK if properly de-identified
2. **Unique combinations:** Count how many individuals share the same values
   - If any combination appears only once or twice → RISK of identification
   - Run frequency counts on combinations of indirect identifiers
3. **Local knowledge:** Could someone with community knowledge identify individuals?
   - Program staff, community members, local officials often know participants
   - Consider who might have access and what they know
4. **Linkage risk:** Could this data be linked with other available data?
   - Public voter rolls, school registries, program participation lists

- Social media, news articles, public records

Small Sample Sizes: Special Considerations

When small samples increase risk:

Data from small, specific populations is often identifiable even when de-identified, and should typically be classified as CONFIDENTIAL at minimum:

Examples requiring caution:

- Interviews with 8 district health officials
- Survey of 12 NGO directors in a region
- Case studies of 5 implementing partners
- Focus groups with 6-8 participants per group
- Program serving 20 beneficiaries in a village

Guidelines for small samples:

Sample Size	Classification Guidance	Recommended Approach with AI
N < 10	Assume CONFIDENTIAL minimum, even if anonymized. Individuals may be identifiable by context.	Do not share verbatim quotes or detailed characteristics with AI without approval. Use local LLMs or consult MIST.
N = 10-50	Likely CONFIDENTIAL if specific population (e.g., all chiefs in district). May be INTERNAL if general population with strong de-identification.	Assess re-identification risk. Remove or generalize characteristics that create unique combinations.
N > 50	Can potentially be INTERNAL if properly de-identified and no other restrictions apply (IRB, donor agreements, etc.).	Ensure k-anonymity (k≥3 recommended). Check that no individual is unique on any combination of variables.

Even with larger samples, be cautious if:

- Population has unique characteristics (refugees in specific camp, beneficiaries of specific program)
- Data includes rare outcomes or characteristics
- Geographic area is small or specific
- Data is longitudinal (multiple observations per person increase identification risk)

Aggregated and Summary Statistics

Can I share summary statistics with AI tools?

Summary statistics generally pose lower risk than individual-level data, but classification depends on several factors:

Likely PUBLIC or INTERNAL (if from approved published sources):

- Means, medians, percentages from published papers
- Summary tables from publicly released reports
- Aggregated statistics with large denominators (n > 50)

Likely INTERNAL (if unpublished but non-sensitive):

- Summary statistics from internal analyses, no PII involved
- Descriptive stats where cell sizes are all n > 10
- Aggregate data that doesn't reveal individual responses

Still CONFIDENTIAL if:

- Unpublished research data under embargo
- Small cell sizes (n < 10) that could reveal individual responses
- Covered by donor agreements restricting data sharing
- Contains sensitive topics even in aggregate (HIV prevalence in small clinic, abuse rates in specific school)

HIGHLY CONFIDENTIAL if:

- Summary statistics of SPI where small numbers could identify individuals
- Could be used to infer individual status (e.g., "2 out of 3 participants tested positive")

Best practices when sharing summary statistics with AI:

- Always suppress cells with n < 10
- For sensitive topics, consider requiring n ≥ 30 for reporting
- Include only statistics, not the underlying data
- If unsure whether data is embargoed, treat as CONFIDENTIAL

Children's Data: Special Rules

Does "children's data" mean ANY data from minors?

Data from or about children requires extra protection, but not all children's data is automatically HIGHLY CONFIDENTIAL:

HIGHLY CONFIDENTIAL (stricter standards):

- Any SPI about children (health, abuse, criminal involvement, disabilities)
- Data from children under 13 in most contexts
- Identifiable photos, videos, or audio of children
- Data that could endanger children if disclosed

- Children in vulnerable situations (refugees, victims of violence, child laborers)

CONFIDENTIAL (standard PII protection):

- Basic demographic data from adolescents (ages 13-17) in low-risk research
- School enrollment data without sensitive information
- Program participation data without health or sensitive topics
- De-identified survey responses from teenagers on non-sensitive topics

Key principle: Apply the "best interest of the child" standard. If disclosure could harm the child, classify as HIGHLY CONFIDENTIAL.

Additional considerations:

- Parental consent requirements are separate from data classification
- Local laws may have different age thresholds (e.g., GDPR uses 16)
- Combined data (child + parent) requires protection at the highest level of either

Practical Decision Tool: Is My Data Properly De-identified?

Use this checklist to assess whether your data is adequately de-identified to reduce classification level:

Level 1: Direct Identifiers (Must ALL be removed)

- ☐ Names, initials, nicknames
- ☐ Contact information (email, phone, address)
- ☐ ID numbers (national ID, passport, participant ID that can be linked)
- ☐ Photos, videos, or audio with identifiable features

Level 2: Indirect Identifiers (Assess and mitigate)

- ☐ Exact ages → Convert to age ranges
- ☐ Specific locations → Aggregate to region or district
- ☐ Exact dates → Round to month or quarter
- ☐ Rare occupations or characteristics → Generalize or suppress
- ☐ Unique combinations → Check frequency counts

Level 3: Re-identification Risk Assessment

- ☐ Sample size > 50 with good diversity?
- ☐ No cells or combinations with n < 3?
- ☐ Community members couldn't identify participants?
- ☐ Data couldn't be linked with public records?

Level 4: Other Restrictions

- ☐ Not subject to IRB restrictions on sharing?
- ☐ Not covered by donor agreement confidentiality clauses?
- ☐ Not embargoed or awaiting publication approval?
- ☐ Doesn't contain information that could harm IPA if disclosed?

If you can check ALL boxes above: Your data may qualify for INTERNAL classification (as long as there is no SPI) for use with IPA-approved AI tools.

If you cannot check ALL boxes: Your data remains CONFIDENTIAL or HIGHLY CONFIDENTIAL. Use only approved tools (Azure OpenAI, local LLMs) or seek approval.

When in doubt: Contact GRDS (researchsupport@poverty-action.org) for guidance on de-identification or MIST for classification questions.

❌ What You CANNOT Share with AI Tools

Regardless of the AI tool you're using, **NEVER** input the following types of information without prior specific use-case approval:

Prohibited

Personally Identifiable Information (PII) from research participants

Unless specifically approved in an IRB protocol (contact GRDS if you need guidance on removing PII from research data, researchsupport@poverty-action.org)

Including but not exclusive to:

- Names, addresses, phone numbers, email addresses
- ID numbers, passport numbers, national ID numbers
- Location data that could identify individuals
- Photos or videos of identifiable people
- Audio recordings with identifiable voices

Sensitive Personal Information (SPI)

Including but not exclusive to:

- Health or medical data
- Biometric data (fingerprints, facial recognition data, DNA)
- Sexual orientation or gender identity information
- Political opinions, religious beliefs, or philosophical views
- Trade union membership
- Criminal history or allegations

IPA Credentials and Access Information

Including but not exclusive to:

- Passwords, API keys, authentication tokens
- System access credentials
- Encryption keys

Information Prohibited by Contract or Legal Agreement

For example:

- Data that donors explicitly restrict from sharing with third parties
- Information covered by non-disclosure agreements (NDAs)
- Proprietary information belonging to research partners, implementers, donors, or other stakeholders
- Intellectual property owned by donors or partners
- Embargoed research findings or data with publication restrictions
- Information shared with IPA under any confidentiality restriction or expectation
- Information restricted by data protection laws or other legal statutes

HR and Legal Conversations

Including but not exclusive to:

- Disciplinary matters
- Performance reviews containing sensitive feedback
- Legal consultations or privileged communications

⚠️ Use Caution With

- **Unpublished Research Data** even if de-identified, may be confidential
- **Embargoed Research Findings** subject to publication restrictions or donor approval requirements
- **Grant Proposals in Development** contain strategic information and budget details
- **Internal Strategy Documents** competitive or sensitive organizational planning
- **Financial Information** budgets, salary data, cost structures
- **Donor or Partner Communications** may contain confidential negotiations, terms, or proprietary information
- **Information from External Parties** if shared with IPA under confidentiality expectations, even informal ones
- **Data Subject to Legal Restrictions**, check grant agreements, data use agreements, and local laws before using AI

> ⓘ Important
>
> If in doubt consult the Legal team or MIST security team to verify the classification of the data and whether it can be shared with AI tools.

Pre-Approved AI Uses

Most AI uses at IPA **do not require special approval.** The following are pre-approved based on data classification:

✅ For PUBLIC Data (No Restrictions)

Any AI tool can be used

- Draft blog posts or communications using published research
- Create summaries of public IPA materials
- Translate published documents
- Generate ideas based on publicly available information
- Literature reviews using published academic papers

✅ For INTERNAL Data (IPA-Approved Tools Only)

- **Claude** (institutional license) - See separate guidance on institutional vs. free licenses
- **Azure OpenAI** (approved for IPA use on Microsoft Azure cloud storage)
- **Local/Offline LLMs** (downloaded to your device, no data transfer)

Appropriate uses:

- Draft internal documents, meeting agendas, project plans
- Improve writing quality and clarity of non-confidential materials
- Brainstorm ideas for programs or approaches
- Translate non-sensitive, nonconfidential or nonproprietary content
- Create internal training materials
- Summarize non-confidential project documentation
- Generate code for data analysis or automation (non-sensitive data)

✅ For CONFIDENTIAL Data (Strictly Limited)

Pre-approved tools (no additional approval needed):

- **Azure OpenAI via API** - Approved for CONFIDENTIAL data including research use
    - **With PII or human subjects data:** Must be documented in IRB protocol BEFORE use
    - **Without PII (e.g., unpublished research data, embargoed findings, grant proposals):** Can use without IRB approval, but still requires proper data handling
    - **IRB-exempt studies:** Can use Azure OpenAI; exemption doesn't prohibit AI use
- **Local/Offline LLMs** - No data leaves your device
    - Approved for any CONFIDENTIAL data
    - Contact GRDS researchsupport@poverty-action.org if you need assistance setting up local LLMs
    - No IRB amendment needed since data stays on your device

⚠️ **Tools requiring case-by-case approval:**

- **Claude institutional license** with data classified as CONFIDENTIAL by any contracts, agreements, or legal obligations. Check any terms in your grant agreements, data use agreements, or donor contracts. If any doubts, check with the Legal team.
- **Other AI tools** - Submit a New Generative AI Use Case Review Request

🛑 For HIGHLY CONFIDENTIAL Data

Contact MIST helpdesk for security review BEFORE using any AI tool

- Requires strict controls and likely needs multiple approvals (DPO, Legal, IT)
- Very limited approved pathways

Common Questions About CONFIDENTIAL Data and AI

💡 Quick Reference: CONFIDENTIAL Data Scenarios

Scenario 1: I have unpublished survey data with NO participant names or contact info, but it's embargoed by my donor

- **Classification:** CONFIDENTIAL (no PII, but embargoed/restricted)
- **Can I use Azure OpenAI?** YES, without IRB approval
- **Can I use Claude?** Verify with grant agreements and project documents
- **Can I use local LLMs?** YES, no approval needed

Scenario 2: I have interview transcripts WITH participant names

- **Classification:** CONFIDENTIAL or HIGHLY CONFIDENTIAL (depends on interview content)
- **Can I use Azure OpenAI?** YES, but MUST be in your IRB protocol FIRST
- **Can I use Claude?** NO, even institutional license not approved for PII
- **Can I use local LLMs?** YES, but still should be in IRB protocol (ethical best practice)

Scenario 3: I'm writing a grant proposal with confidential budget details and preliminary findings (no participant data)

- **Classification:** CONFIDENTIAL (sensitive organizational info)
- **Can I use Azure OpenAI?** YES, no IRB needed
- **Can I use Claude institutional?** Submit request to MIST (or use Azure OpenAI instead)
- **Can I use local LLMs?** YES

Scenario 4: I have de-identified survey data but from a study with only 15 participants in a small village

- **Classification:** CONFIDENTIAL at minimum (small sample = re-identification risk)
- **Can I use Azure OpenAI?** Depends - check if your IRB protocol allows third-party processing
- **Can I use local LLMs?** YES, safest option for small samples
- **Should I get IRB approval?** Check your protocol; small samples may require amendment

Scenario 5: My study was IRB-exempt (anonymous surveys, no PII collected)

- **Classification:** Likely INTERNAL if truly no PII; may be CONFIDENTIAL if embargoed
- **Can I use Azure OpenAI?** YES, exemption doesn't prohibit AI
- **Do I need IRB approval?** NO, exempt studies don't require amendments for AI use with non-PII data

Key Principle: IRB approval is about WHAT data you're processing (PII vs. not), not which tool you use. Azure OpenAI's approval is separate from IRB requirements.

Tool-Specific Guidance

Claude (Anthropic)

IPA has vetted Claude and established an institutional licensing agreement with enhanced data protection terms.

Institutional License vs. Free Claude

IPA Institutional Licenses ($10/user/month, billed to IPA):

Required for:

- Working with INTERNAL data (project documentation, internal memos, non-public materials)
- Accessing IPA's curated Team Projects (e.g., grant proposal development, literature review templates)
- Collaborating with IPA colleagues on shared Projects
- Heavy daily use (institutional license has higher usage limits)

Approved for (with proper classification):

- PUBLIC data: Published research, public materials
- INTERNAL data: Non-confidential organizational content
- LIMITED CONFIDENTIAL data: Only after verifying that use is compliant with any contractual or legal restrictions (e.g., grant agreements, donor contracts). Check with Legal if unsure.

Not approved for:

- PII or identifiable research participant data
- CONFIDENTIAL data that is restricted by donor agreements, IRB protocols, or legal obligations without specific approval
- HIGHLY CONFIDENTIAL data (SPI)

Free Claude Accounts:

Allowed for:

- PUBLIC data only (published research, public blog posts, publicly available information)

Not allowed for:

- INTERNAL data (even non-sensitive organizational content)
- CONFIDENTIAL or HIGHLY CONFIDENTIAL data
- Any work requiring IPA's institutional data protection terms

How to get an institutional license:

1. **Submit a request:** Complete the Claude Institutional License Request Form
2. **Justify business need:** Explain your role and why you need institutional access (e.g., grant writing, research coordination, heavy usage)
3. **License assignment:** MIST will provision your license (typically within 1-2 business days)
4. **Cost:** $10/month per user, billed centrally to IPA (not to individual grants)
5. **Access:** You'll receive invitation to IPA's Claude Team workspace with access to shared Projects

CLAUDE LICENSE QUESTIONS AND ANSWERS

Q: I already have a free Claude account. Can I keep using it?

A: Yes, for PUBLIC data only. If you work with INTERNAL data, you must request an institutional license and switch to using that account for IPA work. You can keep your personal account for non-work use with public information.

Q: I'm writing a grant proposal with confidential budget details. Which Claude should I use?

A: You should NOT use Claude (free or institutional) for CONFIDENTIAL data like grant proposals with sensitive budget information. Instead:

- Use **Azure OpenAI** (approved for CONFIDENTIAL data without special request)
- OR submit a specific use case request to MIST for Claude institutional license approval for this purpose
- OR use **local LLMs**

However, if you're drafting a grant proposal using only PUBLIC research findings and INTERNAL (non-confidential) information, the institutional license is appropriate.

Q: What are IPA's curated Team Projects?

A: IPA maintains shared Claude Projects for common tasks:

- Grant proposal templates and best practices
- Literature review workflows
- Research protocol development assistance
- Communications and blog post drafting
- Data analysis code generation

These are only accessible with an institutional license.

Q: Can I use Claude for quick questions while I'm waiting for my institutional license?

A: Yes, but ONLY for PUBLIC data. If your question involves any INTERNAL, CONFIDENTIAL, or HIGHLY CONFIDENTIAL information, wait for your institutional license or use Azure OpenAI instead.

Q: What if I accidentally used free Claude with INTERNAL data?

A: Follow these steps:

1. Stop using the free account for IPA work immediately
2. Request an institutional license right away
3. Contact MIST if you're concerned about data shared with the free account. File an unexpected event with your IRB if the use included PII.
4. In the future, classify your data before using any AI tool

Q: Does institutional license give me unlimited usage?

A: Institutional licenses have significantly higher usage limits than free accounts, but are not unlimited. Heavy usage may occasionally hit rate limits. If this happens regularly, contact MIST.

Data protection differences:

Feature	Institutional License	Free Claude
Data retention	Commercial terms; not used for training	Consumer terms; less favorable retention
Business Associate Agreement	Available for HIPAA compliance	Not available
Team collaboration	IPA Team Projects, shared workspaces	Individual use only
Usage limits	Higher limits for professional use	Lower limits, rate limiting
Approved data types	PUBLIC, INTERNAL (some CONFIDENTIAL with approval)	PUBLIC only
Support	IPA MIST helpdesk + Anthropic support	Community support only

Bottom line: If you regularly work with INTERNAL data or need to collaborate with IPA colleagues, request an institutional license. It's worth the $10/month investment for the data protection and collaboration features.

Perplexity

- Perplexity has not been evaluated for privacy concerns. Use only for Public data
- Data may be used for training (can opt out in settings)
- Usage limits may apply

ChatGPT (OpenAI)

Free tier:

- Approved for Public data
- Data may be used for training (can opt out in settings)
- Usage limits apply

Paid tiers (Plus, Team):

- Better data protection terms (not used for training)
- Higher usage limits
- Access to more advanced models

> For research use with Confidential data
>
> Use Azure OpenAI instead (see below) for better institutional controls.

Azure OpenAI

This is IPA's primary approved platform for CONFIDENTIAL data, including research uses.

Why Azure OpenAI:

- Runs on IPA's Azure cloud infrastructure
- Data stays within IPA's security perimeter
- Commercial terms with strong data protection
- Not used for training OpenAI's models
- Includes GPT models and Whisper (transcription)

How to access:

- Contact GRDS (researchsupport@poverty-action.org) to request Azure OpenAI API access
- API usage is billed to grants
- Available for research, operations, and other approved use cases

Local/Offline LLMs

Running AI models locally on your device is an approved pathway for working with CONFIDENTIAL data because data never leaves your computer. This is the most privacy-preserving option but requires more technical setup.

Why Choose Local LLMs?

Local LLMs are ideal when:

- You're working with CONFIDENTIAL data and want maximum privacy
- You don't have Azure OpenAI access or IRB approval yet
- You work offline or in locations with unreliable internet
- You want to avoid usage limits or recurring costs
- You're comfortable with some technical setup

When cloud-based tools may be better:

- You need the most advanced capabilities (Azure OpenAI models are generally more capable)

- You want easy setup with no technical configuration
- Your device doesn't meet hardware requirements
- You need real-time collaboration features
- Time is critical and you can't invest in learning local tools

Getting Started with Local LLMs

1. Choose a platform:

Platform	Best For	Ease of Use	Operating Systems
Jan.ai	Non-technical users, graphical interface	⭐⭐⭐⭐⭐ Easy	Windows, Mac, Linux
LM Studio	Desktop users wanting simple setup	⭐⭐⭐⭐⭐ Easy	Windows, Mac, Linux
Ollama	Technical users, command-line interface	⭐⭐⭐ Moderate	Windows, Mac, Linux
GPT4All	Beginners, lightweight option	⭐⭐⭐⭐ Easy	Windows, Mac, Linux

Recommendation for most IPA staff: Start with **Jan.ai** or **LM Studio** for user-friendly graphical interfaces.

2. Check hardware requirements:

Your device needs sufficient RAM and ideally a GPU for good performance:

Model Size	Minimum RAM	Recommended RAM	GPU	Performance Level	Best Use Cases
Small (3B-7B params)	8 GB	16 GB	Optional	Fast but less capable	Quick drafts, brainstorming, simple coding
Medium (7B-13B params)	16 GB	32 GB	Helpful	Good balance	Most research tasks, writing, analysis
Large (13B-34B params)	32 GB	64 GB	Strongly recommended	Best quality	Complex analysis, critical documents

What's a "B parameter" model? The number of parameters (in billions) indicates model size and capability. Larger models are generally more capable but require more resources.

GPU considerations:

- NVIDIA GPUs significantly speed up inference (10-50x faster)
- Apple Silicon Macs (M1/M2/M3) have excellent built-in acceleration
- Models can run on CPU-only, but will be slower
- If you have a gaming laptop or workstation with NVIDIA GPU, you're in great shape

3. Choose a model:

Open-source models vary in size and capability. Start with smaller models to test, then scale up if needed. Models are constantly evolving, so check platforms for latest options. Ollama provides a convenient, searchable database of models: https://ollama.com/search. Huggingface also provides a Model leaderboard

Local vs. Cloud Decision Framework

Use this framework to decide whether to use local LLMs or cloud-based Azure OpenAI:

ⓘ Local vs. Cloud: Decision Guide

Choose LOCAL LLMs when:

- Data is CONFIDENTIAL and you want absolute privacy guarantee
- You have small sample data with re-identification risk
- You're working offline or with intermittent internet
- You want to avoid any data transmission risk
- You have time to set up and learn the tools
- Your hardware meets requirements (16GB+ RAM recommended)
- IRB protocol doesn't explicitly prohibit local processing
- Quality requirements are moderate (drafts, brainstorming, initial analysis)

Choose AZURE OpenAI when:

- You need best-in-class quality (critical documents, publications)
- Speed is important (cloud models respond faster than CPU-only local)
- You're already set up with Azure OpenAI access
- Your device doesn't meet local LLM requirements
- You want zero technical setup time
- You're collaborating with team members using same tool
- Data is CONFIDENTIAL without PII, or you have IRB approval for PII
- Your IRB protocol explicitly approves Azure OpenAI use

Getting Help with Local LLMs

- **GRDS** (researchsupport@poverty-action.org): Can provide guidance on setup, model selection, and research use cases
- **MIST helpdesk**: For device hardware questions or security concerns
- **Peer support**: Ask colleagues who've successfully set up local LLMs. Use the We-Talk-AI channel on Microsoft Teams to connect with others.

AI Meeting Assistants and Embedded AI Features

AI Meeting Notetakers

AI meeting assistants (Otter.ai, Fireflies.ai, etc.) are powerful tools but require careful consideration of consent and data sensitivity.

> ⓘ Consent Requirements
>
> **You MUST obtain explicit consent from ALL participants before using AI notetakers in any meeting**

For IPA Internal Meetings:

- Microsoft Teams has built-in transcription and recording. This feature is only available as a Premium feature. IPA has a limited number of users with this feature enabled. Contact MIST if you have a business need for the feature to be enabled on your account. Transcripts, AI recaps, video recordings, and meeting analytics can be shared with other meeting participants. Note that for recurring meetings these artifacts may be available to people who join the meeting on a different date even if they were not present for the recording.
- Most note-takers will notify meeting participants before transcribing, translating, or recording the meeting. Ask at the start: "I'd like to use [AI notetaker] to transcribe this meeting. Is everyone comfortable with that?"
- If *anyone* declines, do not use the tool
- Consider the content: if discussing Confidential information, do not use AI notetakers unless specifically approved

For External Meetings (with donors, partners, research participants):

- Always ask permission of all participants before recording or transcribing
- Explain what the tool does and how data will be used
- Respect any decline without pressure
- Consider whether the meeting content is appropriate for AI processing

For HR or Legal Conversations:

- **AI notetakers are PROHIBITED** in conversations with HR or Legal
- These conversations are privileged and confidential

When Someone Else Brings an AI Notetaker to Your Meeting: You have the right to decline:

- "I'm not comfortable with AI recording this conversation. Could we turn that off?"
- "IPA policy requires consent for AI transcription. I need to decline."
- If they insist, you can reschedule or escalate to your supervisor

Embedded AI Features in Other Software

Many tools now include built-in AI features. **The same data classification principles apply.**

Examples of embedded AI:

- **Microsoft Copilot** (in Word, Excel, PowerPoint, Teams, Outlook)
- **Google Workspace AI** (in Docs, Sheets, Gmail)
- **Grammarly** (writing assistant)
- **Zoom AI Companion** (meeting summaries, chat assistance)
- **Notion AI**
- **Slack AI**
- **GitHub Copilot** (coding assistant, available through IPA's GitHub account, license billed to grants)

Before using embedded AI features:

1. **Classify your data**: What are you working with?
2. **Understand data flow**: Does the AI feature send your content to the vendor's servers? Or does it process locally?
3. **Check IPA approval**:
    - For **Internal data or less sensitive**: If the main tool is IPA-approved (Microsoft 365, Google Workspace), AI features are generally okay
    - For **Confidential data**: Verify with MIST that the specific AI feature has been reviewed
4. **Apply common sense**:
    - Don't use Grammarly on grant proposals with Confidential budget information unless approved
    - Don't let Zoom AI summarize meetings discussing sensitive research data
    - Don't use AI features in tools that haven't been reviewed by IPA IT

> **Tip**
>
> When in doubt, contact MIST helpdesk to verify approval status.

Research-Specific Guidance

As a research organization, IPA staff frequently work with sensitive data. Here's specific guidance for using AI in research contexts:

IRB Protocols and AI Use

When IRB Approval Is Required

You MUST disclose AI use in your IRB protocol when:

AI will process data that contains:

- PII (names, contact information, identifiers)
- Identifiable information (indirect identifiers in small samples)
- Any data where participants could be re-identified

AND the data was collected from human subjects under IRB oversight

AI uses requiring IRB disclosure:

- Transcribing interviews containing participant identifiers
- Analyzing qualitative data with names, locations, or identifiable details
- Coding interview transcripts before they've been de-identified
- Processing survey data that includes PII fields
- Using AI on focus group recordings or transcripts with participant identifiers

AI uses NOT requiring IRB disclosure:

- Processing fully de-identified research data (all identifiers removed, low re-identification risk)
- Analyzing published or publicly available data
- Using AI to draft research protocols, consent forms, or survey instruments
- Literature reviews and synthesis of published papers
- Coding or analysis of data that was never collected under IRB (e.g., secondary data analysis of anonymous public datasets)
- IRB-exempt studies using AI on the anonymous data collected

What to include in your IRB protocol or amendment:

Required Element	What to Specify	Example Language
AI Tool	Specific platform and model	"Azure OpenAI (GPT-4) via IPA's secure API" or "OpenAI Whisper for transcription"
Data Types	What data AI will process	"Audio recordings of participant interviews will be transcribed" or "De-identified interview transcripts will be coded for themes"
Data Protection	Security measures during AI processing	"Data transmitted via encrypted API; processed on IPA's Azure infrastructure; not used for model training"
Data Transmission	Where data goes	"Data sent to Azure OpenAI cloud service (US region)" or "Local processing on researcher device, no external transmission"
Privacy Measures	How you'll protect participants	"Recordings will be deleted after transcription; transcripts will be de-identified before analysis; no participant names will be included in prompts"

Data Protection Assessments (DPAs)

For some uses of Confidential research data that don't involve PII, you may need a **Data Protection Assessment (DPA)**. DPAs help identify and mitigate risks related to:

- Research ethics
- Donor or partner data ownership

When might you need a DPA?

- Using AI to analyze sensitive but de-identified research data
- Processing data where donors have specific requirements
- Using new AI tools not yet reviewed by IPA

How to request a DPA: Submit a request to MIST helpdesk (see "How to Request Approval" section below)

Common Research Scenarios

Scenario 1: Literature Review

- **What**: Using AI to find and summarize published academic papers
- **Classification**: Public data
- **Approval needed?** No, use any AI tool
- **Recommended tools**: Claude (with Projects for shared searches), ChatGPT, or regular web-based AI

Scenario 2: Transcribing Interviews (with PII)

- **What**: Converting audio recordings of participant interviews to text
- **Classification**: Highly Confidential (contains PII and sensitive information)
- **Approval needed?** Yes - must be in IRB protocol
- **Recommended tool**: Azure Whisper (OpenAI's transcription service on Azure)
- **Alternative**: Local transcription using offline Whisper model

Scenario 3: Coding De-identified Interview Transcripts

- **What**: Using AI to help identify themes in qualitative data that has been properly de-identified
- **Classification**: Confidential (research data, but no PII)
- **Approval needed?** Depends - if using Azure OpenAI or local LLMs, no new approval needed. If using other tools, submit request to MIST
- **Recommended tools**: Azure OpenAI, local LLMs

Scenario 4: Analyzing Survey Data (Quantitative)

- **What**: Using AI to help write code for statistical analysis, interpret results
- **Classification**: Depends on whether dataset contains PII
- **Approval needed?** If no PII in the dataset shared with AI, likely no approval needed for Internal-classified data

- **Note**: Be cautious about sharing complete datasets; consider sharing just summary statistics or code

Scenario 5: Draft Research Protocol or Survey Instrument

- **What**: Using AI to help write research protocols, questionnaires, or measurement instruments
- **Classification**: Internal (if no PII or confidential study design)
- **Approval needed?** No - use IPA-approved tools
- **Recommended**: Claude Projects has tools for this; ChatGPT also works well

Scenario 6: Grant Proposal with Confidential Budget Details

- **What**: Using AI to help write a proposal that includes sensitive budget information or unpublished research findings
- **Classification**: Confidential
- **Approval needed?** Use IPA-approved tools; Claude institutional license recommended for collaboration
- **Note**: Be cautious about including specific budget numbers or proprietary methodologies

When You Need Approval for New Use Cases

You need to submit a request for approval when:

- **Using AI tools with CONFIDENTIAL or HIGHLY CONFIDENTIAL data that isn't covered by existing approvals**
- **Using an AI tool that isn't on the approved list** (Claude, ChatGPT, Azure OpenAI, local LLMs)
- **Proposing a new type of AI use** that hasn't been done at IPA before
- **Working in a country with strict data protection laws** and unsure if your use case is compliant
- **Processing data governed by specific donor agreements** that may restrict AI use

Examples Requiring Approval

- "I want to use Claude to analyze interview transcripts that contain indirect identifiers"
- "Can I use {new AI tool X} to process our research data?"
- "I want to use AI to help with recruitment screening and candidate evaluation"
- "Can we use AI to translate beneficiary surveys in {specific country}?"

Examples NOT Requiring Approval

- "I want to use Claude to help me write a blog post about our published research"
- "Can I use ChatGPT to draft internal meeting agendas?"
- "I want to use Azure OpenAI to transcribe interviews (and it's in my IRB protocol)"
- "Can I use Claude's literature review Project to find relevant papers?"

How to Request Approval

Single Point of Contact: MIST Helpdesk

All requests for new AI use case approvals should go through the MIST helpdesk:

Submit a ticket: https://ipahelpdesk.freshservice.com/support/tickets/new

Subject line MUST include: "AI Use Case"

What to include in your request:

1. **Data Classification**: Classify the data you plan to use according to the table in this document
    - Public
    - Internal
    - Confidential
    - Highly Confidential
2. **Description of Use**:
    - What do you want to do with the AI tool?
    - What specific tool or service do you want to use?
    - Why is this tool/approach beneficial for your work?
3. **Data Details**:
    - What types of information will the AI process?
    - Does it include PII? If so, what types?
    - Is it governed by an IRB protocol, donor agreement, or other restrictions?
    - Where is the data currently stored?
4. **Context**:
    - What is your role/project?
    - Is this for research, business development, operations, or communications?
    - What country/countries does this involve?
5. **Timeline**: When do you need approval by?

What Happens Next:

MIST will route your request to the appropriate reviewers based on the nature of your use case:

- Legal (for contractual/IP issues)
- Data Protection Officer (for privacy/GDPR issues)
- IT Security (for technical security assessment)

- Research team (for IRB-related questions)

You may be asked for additional information or clarification. Please build in time for review and approval.

Pro Tip: For complex use cases, consider scheduling a consultation with MIST, Legal, or the DPO BEFORE submitting your formal request. Early consultation can help you design your use case appropriately and speed up approval.

Using AI Responsibly: Quality and Limitations

AI tools are powerful but have significant limitations. Never trust AI output without verification.

Hallucination: AI Makes Things Up AI models generate content that sounds authoritative but may be completely fabricated.

What Gets Fabricated:

- Citations to papers that don't exist
- Statistics and facts with no real source
- Quotes attributed to people who never said them
- Technical or legal information that's incorrect
- Answers to logical or mathematical problems

Why: AI predicts plausible-sounding text, not truth. It "hallucinates" reasonable-seeming content even when no source supports it.

What You Must Do:

- Verify every factual claim through trusted sources
- Check every citation - look up papers directly (you're encouraged to explicitly require a DOI or URL for any references that the AI tool identifies)
- Cross-reference all statistics with original data
- Treat AI as a draft requiring fact-checking
- Do not use AI as a calculator unless you ask it to show you each step

Never use unverified AI output in: grant applications, published research, legal/regulatory submissions, or anywhere accuracy is critical.

Bias: AI Reflects Training Data Biases AI perpetuates biases from its training data, including gender, racial, geographic, socioeconomic, and cultural biases.

Examples that arise in our work:

- Suggesting interventions inappropriate for low-income settings
- Using deficit language about communities we serve
- Making cultural assumptions that don't apply locally
- Overlooking power dynamics in research relationships

What You Must Do:

- **Review output** critically for biased language
- **Apply your expertise** about local contexts and cultures
- **Edit content** that perpetuates stereotypes
- **Use AI as a starting point**, then apply human judgment

Reproducibility: Same Question ≠ Same Answer

AI generates different outputs for the same input due to probabilistic generation.

Implications:

- Cannot rely on AI for reproducible research analysis
- Same prompt may yield different results across users or time
- Not suitable for automated decisions without human review

What You Must Do:

- Save prompts and outputs if documentation is needed
- Don't rely on AI for analysis requiring exact reproducibility
- Verify critical outputs independently
- Document which version/model you used

Verification Requirements For ALL AI-assisted work:

- Fact-check: Verify all claims using authoritative primary sources
- Apply expertise: AI doesn't replace your professional judgment
- Edit thoroughly: AI output is a starting point, not finished work
- Extra care for: research findings, grant proposals, communications about participants, legal/compliance content, published materials

You remain responsible for accuracy and quality of your work, even when AI helped create it.

Disclosing AI Use

Transparency about AI use is important in specific contexts, but most routine AI use doesn't require disclosure.

When Disclosure Is Required

External requirements:

- **Journal policies:** Many academic journals now require AI disclosure - check submission guidelines
- **Donor/funder requirements:** Some RFPs or grant agreements specify AI disclosure - check your agreement
- **IRB protocols:** If AI processes data containing PII or human subjects data, include in your protocol
- **When asked:** If donors, partners, or colleagues ask how content was produced, be transparent

Substantive intellectual work by AI:

- **Summarization/synthesis:** If AI analyzed or synthesized content that readers would assume a human reviewed (e.g., "we analyzed 200 open-ended interviews and identified three themes" when AI did the initial synthesis)
- **Data analysis:** If AI performed analysis that's presented as findings (especially in research contexts)
- **Content where AI did the thinking:** If AI generated substantive conclusions, interpretations, or recommendations that you're presenting

When Disclosure Is NOT Necessary

You're using AI as a writing tool when you:

- Provide all the content and direction
- Use AI to articulate or improve your own ideas
- Draft with AI but verify everything and apply your expertise
- Use AI like spell-check, grammar tools, or translation

AI as a writing tool doesn't require disclosure, just like you don't disclose using spell-check, search engines, or calculators.

How to Disclose (When Required)

Keep it simple:

- **Which tool:** "Claude," "ChatGPT," "Azure OpenAI"
- **What it did:** "assisted with literature synthesis," "transcribed interviews," "analyzed qualitative themes"

- **What you verified:** “all content reviewed and verified by author”

For example, *this AI usage guidelines document* was produced with the help of Claude Sonnet 4.5 (as evidenced by the distinctive formatting). All content was revised extensively by the author with input from several humans.

Frequently Asked Questions

Q: How do I know if my donor or grant agreement restricts AI use?

A: Review your grant agreement, data use agreement, or contract for:

- Restrictions on sharing data with third parties
- Requirements for data to stay within specific geographic regions
- Confidentiality clauses or NDAs
- Intellectual property ownership provisions
- Publication or embargo restrictions

If you're unsure whether your agreement permits AI use, contact Legal through MIST helpdesk before proceeding. Err on the side of caution - classify as Confidential and seek approval.

Q: I'm not sure what classification my data falls into. What should I do?

A: When in doubt, choose the more restrictive classification. You can also contact MIST helpdesk or your supervisor for guidance. It's better to be cautious and seek clarification than to misclassify and potentially expose sensitive information.

Q: Can I use AI tools to help me code or write software?

A: Yes! For coding assistance:

- **GitHub Copilot** is available through IPA's GitHub account - use this for coding work
- **Claude, ChatGPT, or local LLMs** can help with code review, debugging, explaining concepts
- **Be cautious** about sharing proprietary code or code that processes Confidential data

Q: What if I'm already using an AI tool and now realize I might need approval?

A: Stop using the tool with that data immediately and submit a request to MIST helpdesk. Be transparent about what you've already done - this helps us assess any risks and determine next steps. You won't be penalized for asking!

Q: Can I use AI to help write my performance review or HR documents?

A: For self-assessments or draft inputs, yes (using IPA-approved tools with Internal classification). However, do not input other people's confidential performance information, disciplinary matters, or HR conversations into AI tools.

Q: Our donor has asked us to use a specific AI tool. Can I use it?

A: Submit a request to MIST helpdesk with “AI Use Case” in the subject. Include information about the donor requirement and what data will be processed. MIST will work with Legal and DPO to review.

Q: What about AI tools embedded in equipment or software we purchase (e.g., survey platforms with AI analysis features)?

A: The same principles apply. Before using AI features in purchased software:

1. Classify the data the software will process
2. If Confidential or Highly Confidential, verify with MIST that the software/feature is approved
3. Ensure any use with PII is covered in your IRB protocol

Q: Can I use AI to analyze social media data or publicly available online content? A: Generally yes, since this is Public data. However, be aware of:

- Ethical considerations around using publicly posted but potentially sensitive content
- Terms of service for the platforms you’re scraping/analyzing
- Whether “public” data might still contain PII that requires protection

Q: I want to experiment with new AI tools to see if they could help my work. How do I do that responsibly?

A: Innovation is encouraged. Start by:

1. Experimenting with PUBLIC data only (published papers, public IPA content, hypothetical examples)
2. Testing with synthetic/fake data that mimics your real data structure
3. Once you understand the tool’s potential, submit a use case request before using it with real Confidential data

Q: How do I know if an AI tool “sees” my data or if it stays private?

A: Look for:

- **Privacy policy**: Does it say data is used for training or analytics?
- **Terms of service**: Are you granting the provider broad rights to use your content?
- **Local vs. cloud**: Does the tool process data on your device or send it to servers?
- **Enterprise agreements**: Does IPA have a commercial agreement with specific data protection terms?

If unclear, contact MIST before using the tool with Confidential data.

Q: What if I see a colleague using AI in a way that seems risky or inappropriate?

A: Please speak up! You can:

1. Raise it directly with the colleague if you're comfortable
2. Discuss with your supervisor
3. Contact MIST helpdesk confidentially
4. Use IPA's whistleblower mechanisms if appropriate

We're all learning together, and peer accountability helps keep everyone safe.

Q: These guidelines seem restrictive. Why all the caution?

A: We're actually trying to be enabling, not restrictive! The data classification approach means that most of your day-to-day AI use (Public and Internal data) doesn't require special approval. The caution is focused on Confidential and Highly Confidential data where we have legal, ethical, and contractual obligations. The goal is protecting research participants, donors, and IPA while enabling innovation.

Q: Will these guidelines be updated as AI technology evolves?

A: Yes! AI is evolving rapidly. These guidelines will be reviewed and updated regularly. If you have suggestions for improvements or encounter situations not covered here, please let MIST know.

Contact and Support

For AI Use Case Approvals: MIST Helpdesk: https://ipahelpdesk.freshservice.com/support/tickets/new (Include “AI Use Case” in subject line)

Additional Resources:

- MIST1133: Artificial Intelligence Usage Policy (formal policy)
- MIST1073: Data Classification Policy (formal policy)
- Understanding IPA’s Claude Institutional Licenses (companion document)

Document Change Log

Version	Date	Changes	Author	Contributors
1.0	9 Mar 2026	Initial release	Steven Glazerman	Niall Keleher, Ernest Ostro